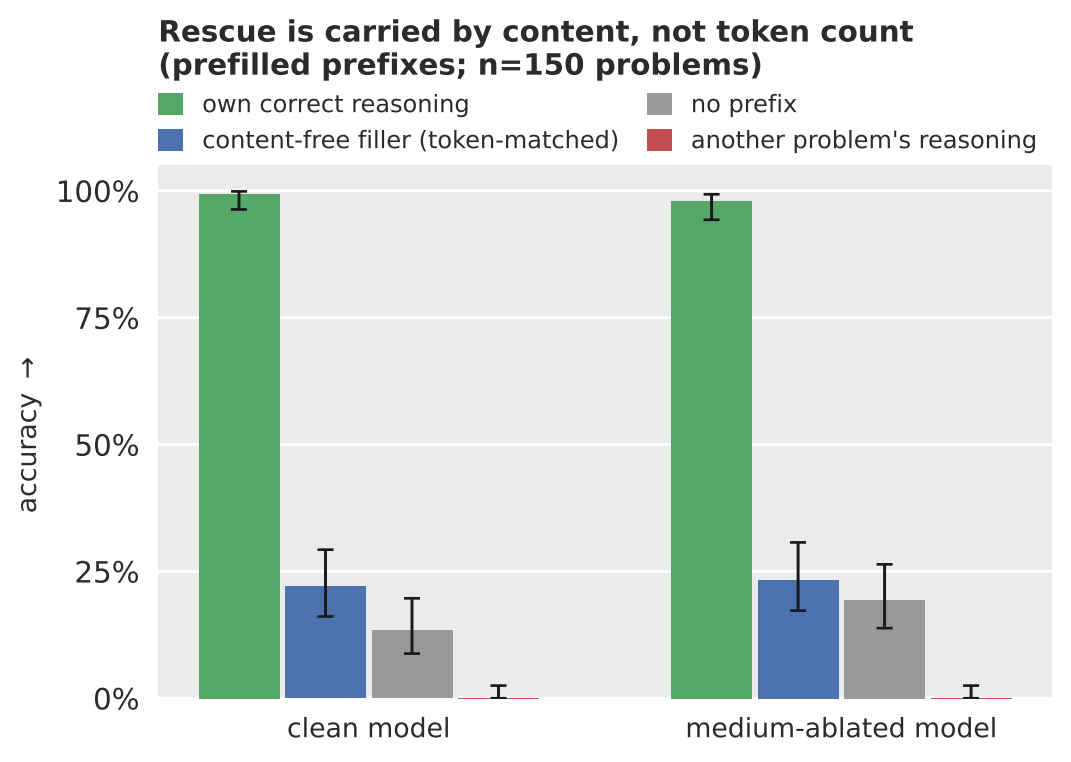
Rescue is carried by content, not token count
(prefilled prefixes; n=150 problems)
own correct reasoning
content-free filler (token-matched)
no prefix
another problem's reasoning
100%
75%
50%
25%
0%
accuracy →
clean model
medium-ablated model